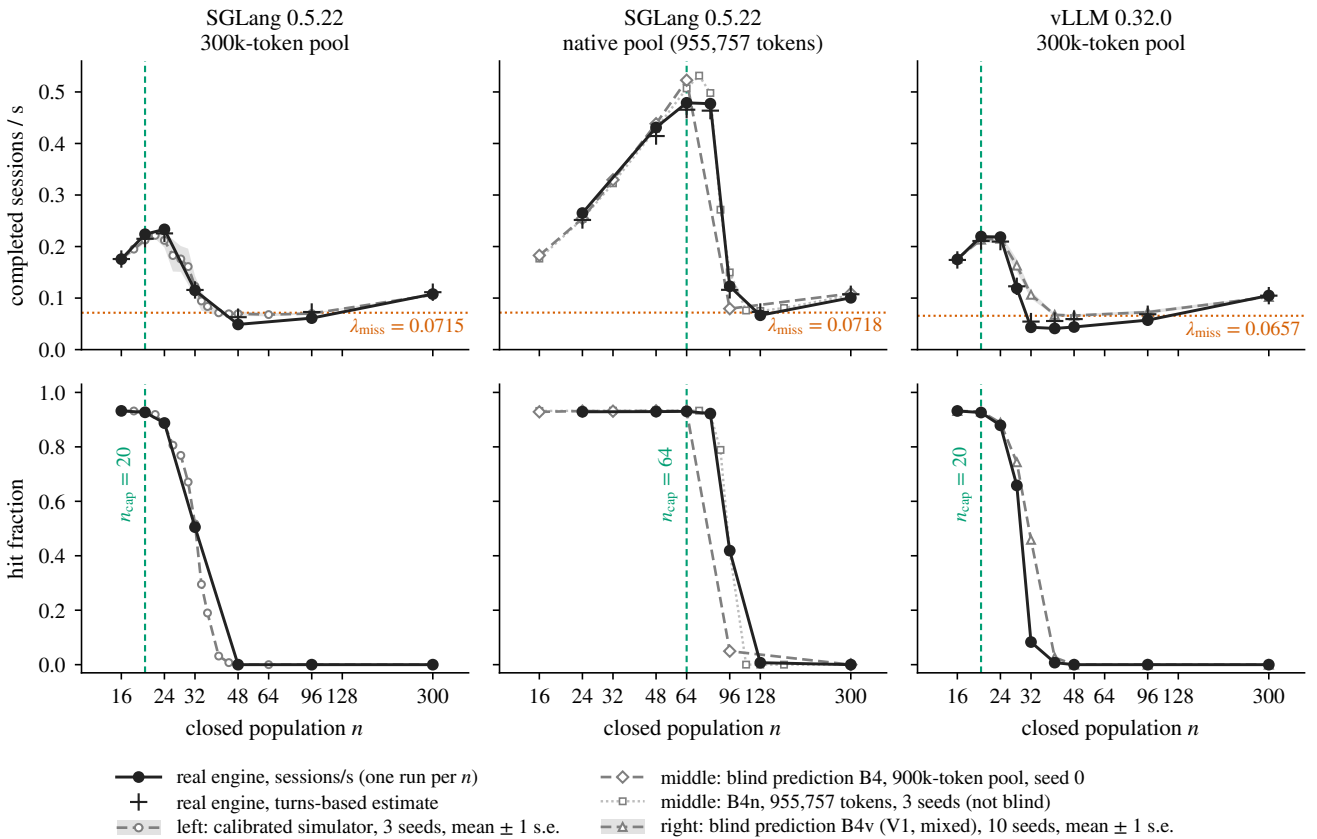
SGLang 0.5.22
300k-token pool
SGLang 0.5.22
native pool (955,757 tokens)
vLLM 0.32.0
300k-token pool
completed sessions / s
hit fraction
$\lambda_{miss} = 0.0715$
$\lambda_{miss} = 0.0718$
$\lambda_{miss} = 0.0657$
$n_{cap} = 20$
$n_{cap} = 64$
$n_{cap} = 20$
closed population n
real engine, sessions/s (one run per n)
real engine, turns-based estimate
left: calibrated simulator, 3 seeds, mean ± 1 s.e.
middle: blind prediction B4, 900k-token pool, seed 0
middle: B4n, 955,757 tokens, 3 seeds (not blind)
right: blind prediction B4v (V1, mixed), 10 seeds, mean ± 1 s.e.