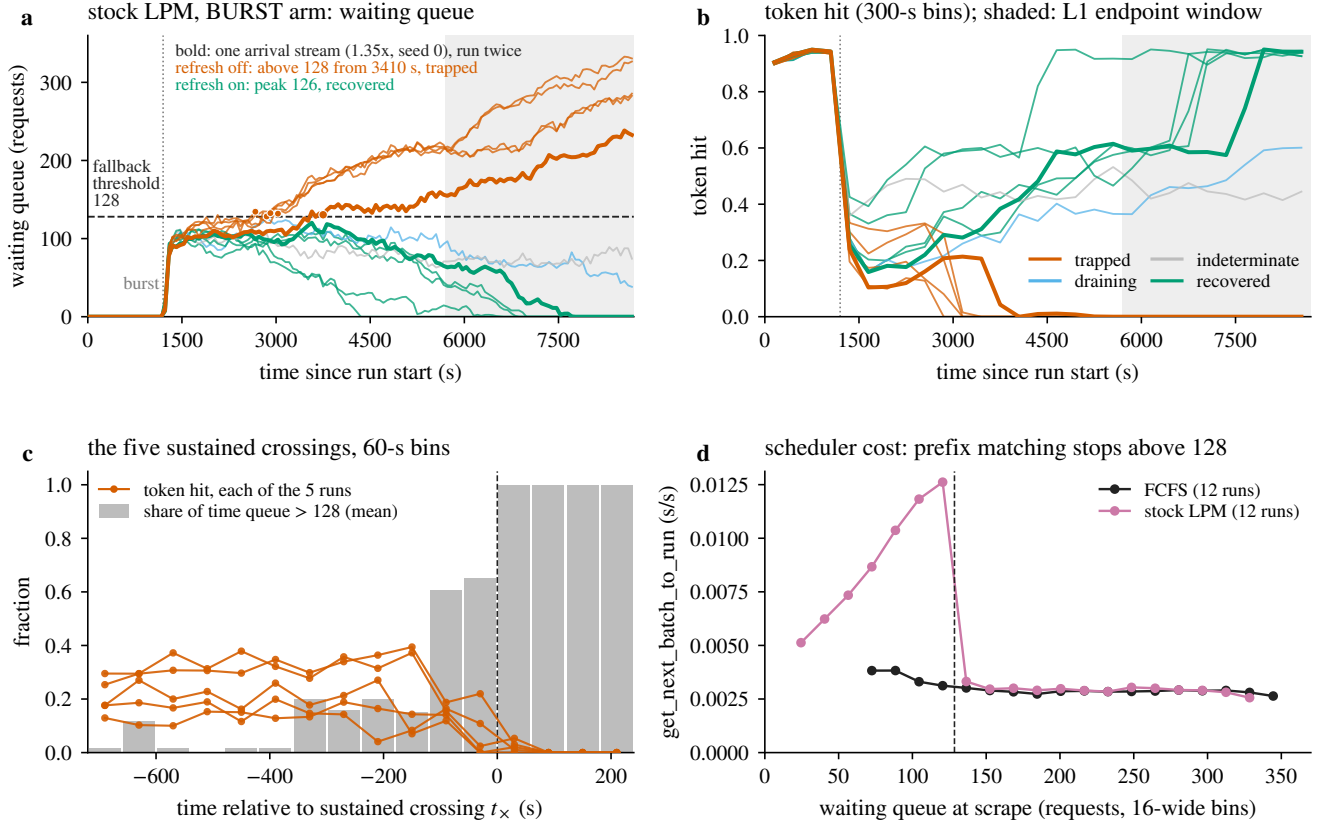
a
stock LPM, BURST arm: waiting queue
bold: one arrival stream (1.35x, seed 0), run twice
refresh off: above 128 from 3410 s, trapped
refresh on: peak 126, recovered
fallback threshold 128
burst
waiting queue (requests)
time since run start (s)
b
token hit (300-s bins); shaded: L1 endpoint window
token hit
trapped
draining
indeterminate
recovered
time since run start (s)
c
the five sustained crossings, 60-s bins
token hit, each of the 5 runs
share of time queue > 128 (mean)
fraction
time relative to sustained crossing $t_\times$ (s)
d
scheduler cost: prefix matching stops above 128
FCFS (12 runs)
stock LPM (12 runs)
get_next_batch_to_run (s/s)
waiting queue at scrape (requests, 16-wide bins)