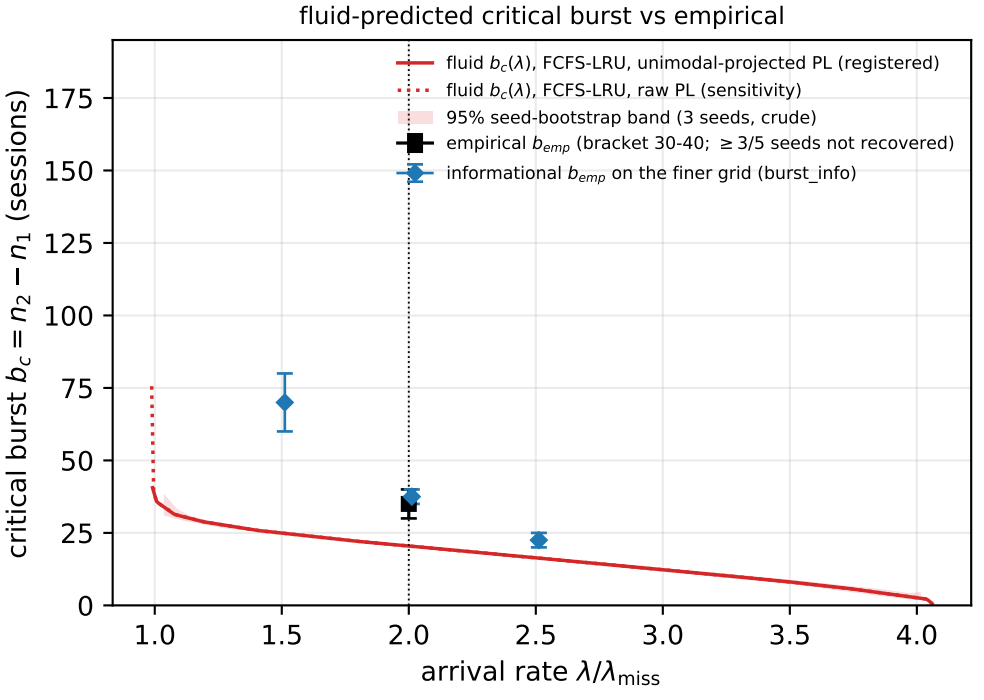
fluid-predicted critical burst vs empirical
fluid $b_c(\lambda)$, FCFS-LRU, unimodal-projected PL (registered)
fluid $b_c(\lambda)$, FCFS-LRU, raw PL (sensitivity)
95% seed-bootstrap band (3 seeds, crude)
empirical b_{emp} (bracket 30-40; ≥ 3/5 seeds not recovered)
informational b_{emp} on the finer grid (burst_info)
critical burst $b_c = n_2 - n_1$ (sessions)
arrival rate $\lambda/\lambda_{\text{miss}}$
0
25
50
75
100
125
150
175
1.0
1.5
2.0
2.5
3.0
3.5
4.0

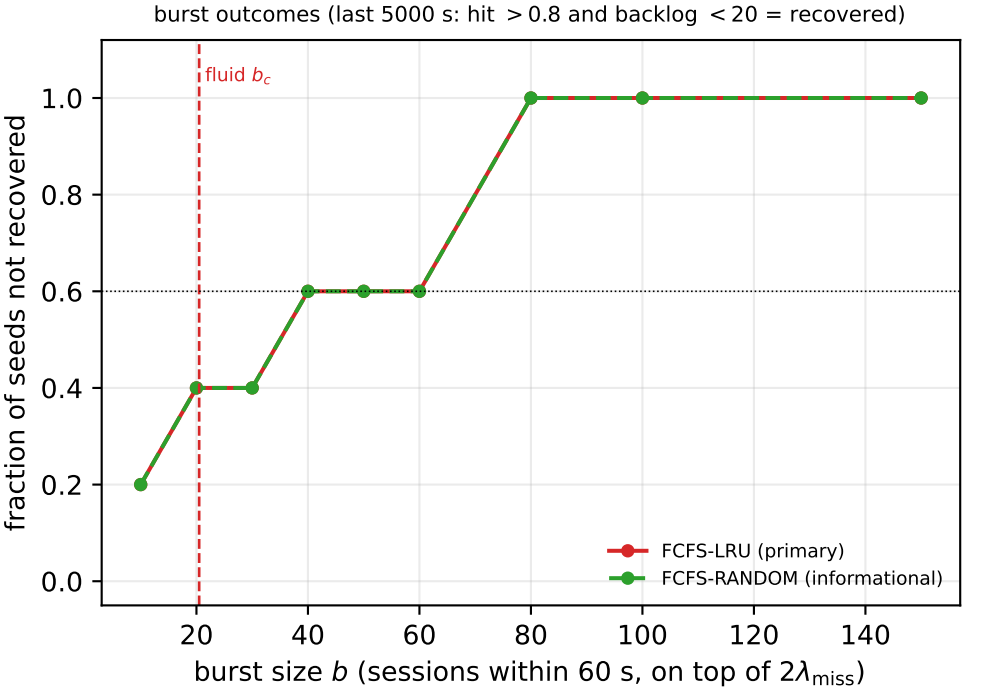
burst outcomes (last 5000 s: hit > 0.8 and backlog < 20 = recovered)
fluid b_c
fraction of seeds not recovered
burst size b (sessions within 60 s, on top of $2\lambda_{\text{miss}}$)
FCFS-LRU (primary)
FCFS-RANDOM (informational)
0.0
0.2
0.4
0.6
0.8
1.0
20
40
60
80
100
120
140